



2017 UNITEC-NTHU Summer School  
on the Frontier of Information Technology

# Deep Learning Lab (Prof. Min Sun) -- TensorFlow Tutorial


Speaker: Tz-Ying (Gina) Wu

# Outline


- Introduction to TensorFlow
- TensorFlow sample codes
  - Linear regression
  - Neural Network (classification)
- Exercise
  - CNN
  - RNN

# Introduction to TensorFlow

# What is TensorFlow?



multidimensional  
data array



computation using data  
flow graphs

# What is TensorFlow?

- TensorFlow is a ***deep learning*** library open-sourced by Google in 2015
- provides primitives for defining functions on tensors and automatically computing their derivatives

You don't need to write  
backpropagation by yourself

- Support CPU-only, GPU usage



# To write a TensorFlow program, we need to ...

- Build a graph (define your model)
- Create a session
- Run the session
  - Initialize variables (if there are variables in the graph)
  - Feed the data
  - Run the graph

# To write a TensorFlow program, we need to ...

- Build a graph (define your model)
- Create a session
- Run the session
  - Initialize variables (if there are variables in the graph)
  - Feed the data
  - Run the graph

# Build a graph (define your model)

- All the computations in TensorFlow graph are tensor operations

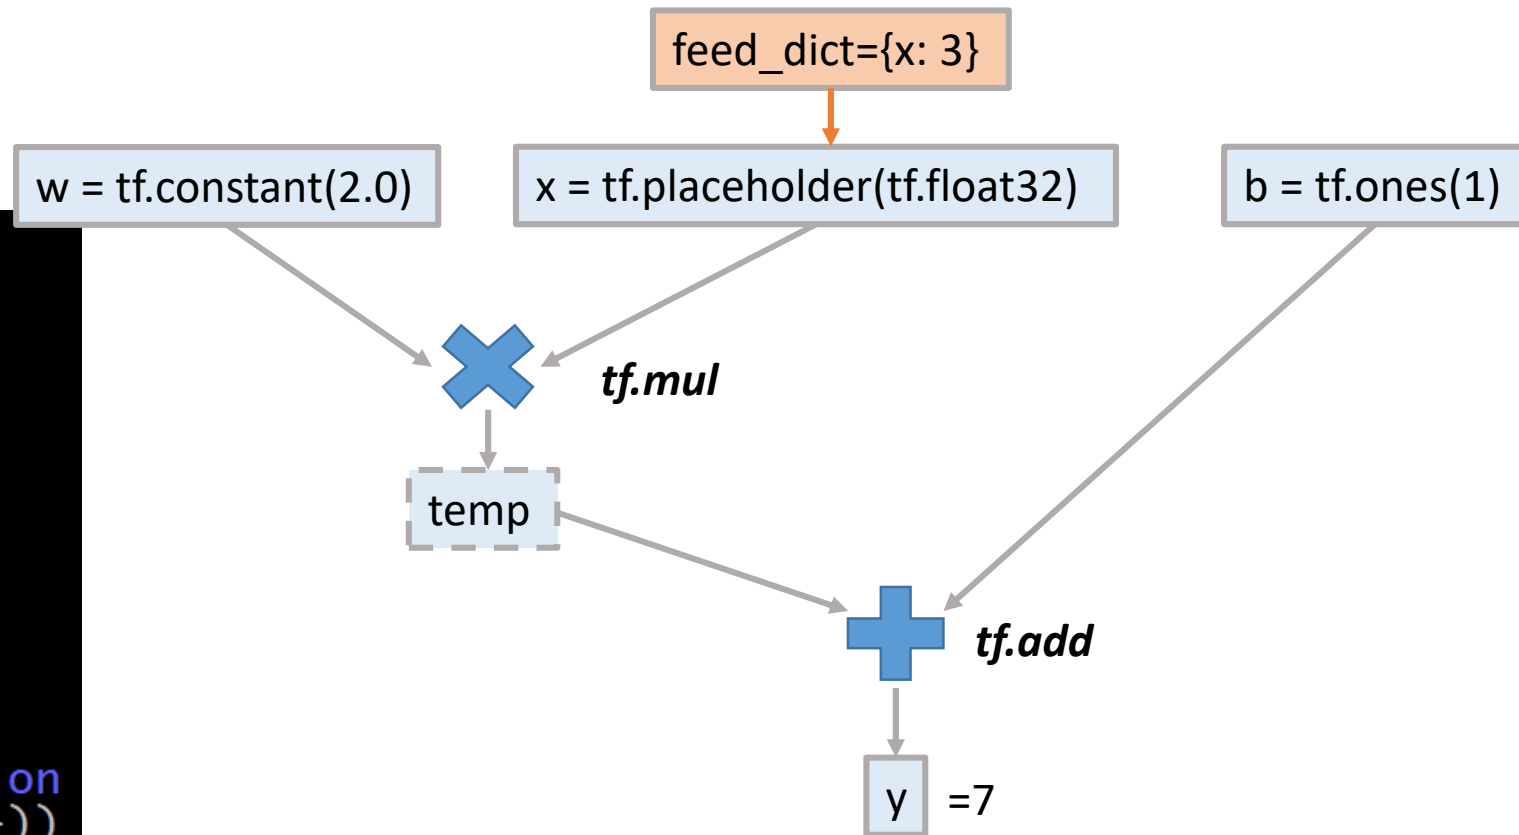
E.g.  $y = w \cdot x + b$

```
import tensorflow as tf

# build the graph
w = tf.constant(2.0)
x = tf.placeholder(tf.float32)
b = tf.ones(1)




# y = w*x+b
y = tf.add(tf.mul(w, x), b)

# create a session
sess = tf.Session()
# feed the data, and run the session
print(sess.run(y, feed_dict={x: 3}))
```





# Build a graph (define your model)

- Tensors can be declared by various ways, e.g.,
  - `tf.zeros((2,2)), tf.ones((1, 2, 3))`  `np.zeros((2, 2), np.ones((1, 2, 3))`
  - `tf.constant([2, 3])`  `np.array([2, 3])`
  - `tf.Variable(tf.zeros((2,2)), name="weights")`  `np.array([[0, 0],[0, 0]])`
  - `tf.placeholder(tf.float32, shape=(10, 1))`
- etc.
- Variables should be initialized before running
- ***tf.placeholder()*** is to reserve the place for input data

# To write a TensorFlow program, we need to ...

- Build a graph (define your model)
- Create a session
- Run the session
  - Initialize variables (if there are variables in the graph)
  - Feed the data
  - Run the graph

# Create a session

- “A **Session** object encapsulates the environment in which **Operation** objects are executed, and **Tensor** objects are evaluated.” - [TensorFlow Docs](#)
- Use ***tf.Session()*** or ***tf.InteractiveSession()*** to create a session

```
sess = tf.Session()  
...  
sess.close()
```

or

```
with tf.Session() as sess:  
    ...
```

# To write a TensorFlow program, we need to ...

- Build a graph (define your model)
- Create a session
- Run the session
  - Initialize variables (if there are variables in the graph)
  - Feed the data
  - Run the graph

# Initialize variables

- “The **Variable()** constructor requires an initial value for the variable, which can be a **Tensor** of any type and shape. ” - [TensorFlow Docs](#)
- Declaration: ***tf.Variable(<initial-value>, name=<optional-name>)***
  - ***<initial-value>*** can be a fixed-value tensor or be random initialized from a distribution
  - E.g. `tf.Variable(tf.zeros((2,2)), name="weights")`
  - E.g. `tf.Variable(tf.random_uniform([100, 2], -1.0, 1.0))`
- Initialization: ***sess.run(tf.initialize\_all\_variables())***
- [Optional: restore parameters from a TensorFlow model (use Saver)]

# Feed the data

- Tensorflow provide ***feed\_dict*** as the bridge between **numpy array** and **tensor**
- Usage: *feed\_dict={<placeholder\_name>: <numpy\_array>}*
- e.g.

declaring placeholder  
when building the graph

***The shape of the placeholder  
and the data fed in must be  
same!!!***

feed the data into the  
placeholder when  
running the session

```
import tensorflow as tf
import numpy as np

# build the graph
w = tf.constant(2.0)
x = tf.placeholder(tf.float32, [2, 2])
b = tf.ones([2, 2])

# y = w*x+b
y = tf.add(tf.mul(w, x), b)

# create a session
sess = tf.Session()
# feed the data, and run the session
x_data = np.array([[1, 2], [3, 4]])
print(sess.run(y, feed_dict={x: x_data}))
```

# Feed the data

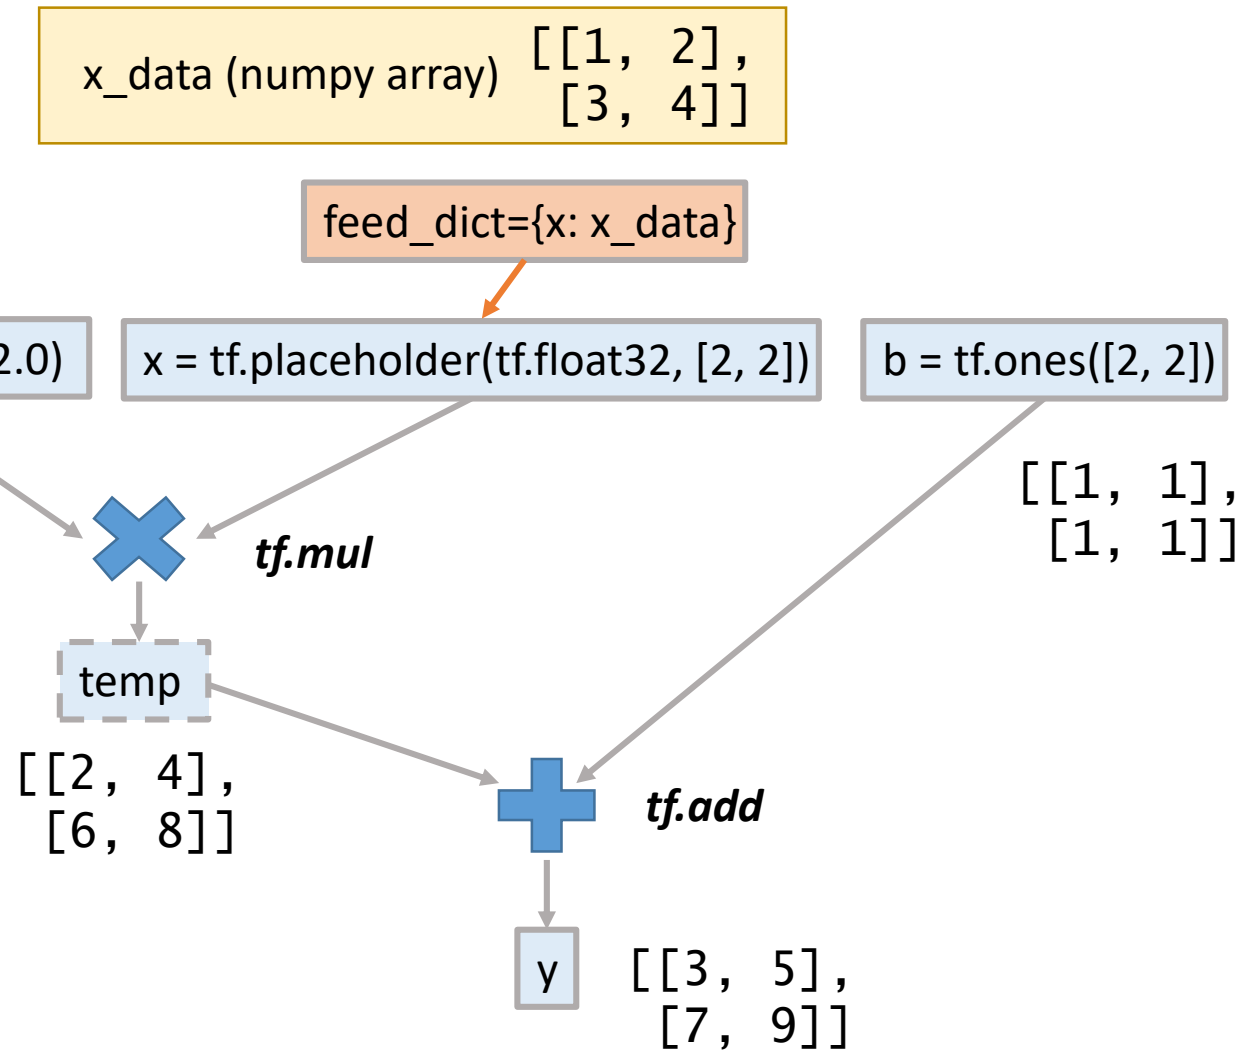
E.g.  $y = w \cdot x + b$

```
import tensorflow as tf
import numpy as np

# build the graph
w = tf.constant(2.0)
x = tf.placeholder(tf.float32, [2, 2])
b = tf.ones([2, 2])

# y = w*x+b
y = tf.add(tf.mul(w, x), b)

# create a session
sess = tf.Session()
# feed the data, and run the session
x_data = np.array([[1, 2], [3, 4]])
print(sess.run(y, feed_dict={x: x_data}))
```



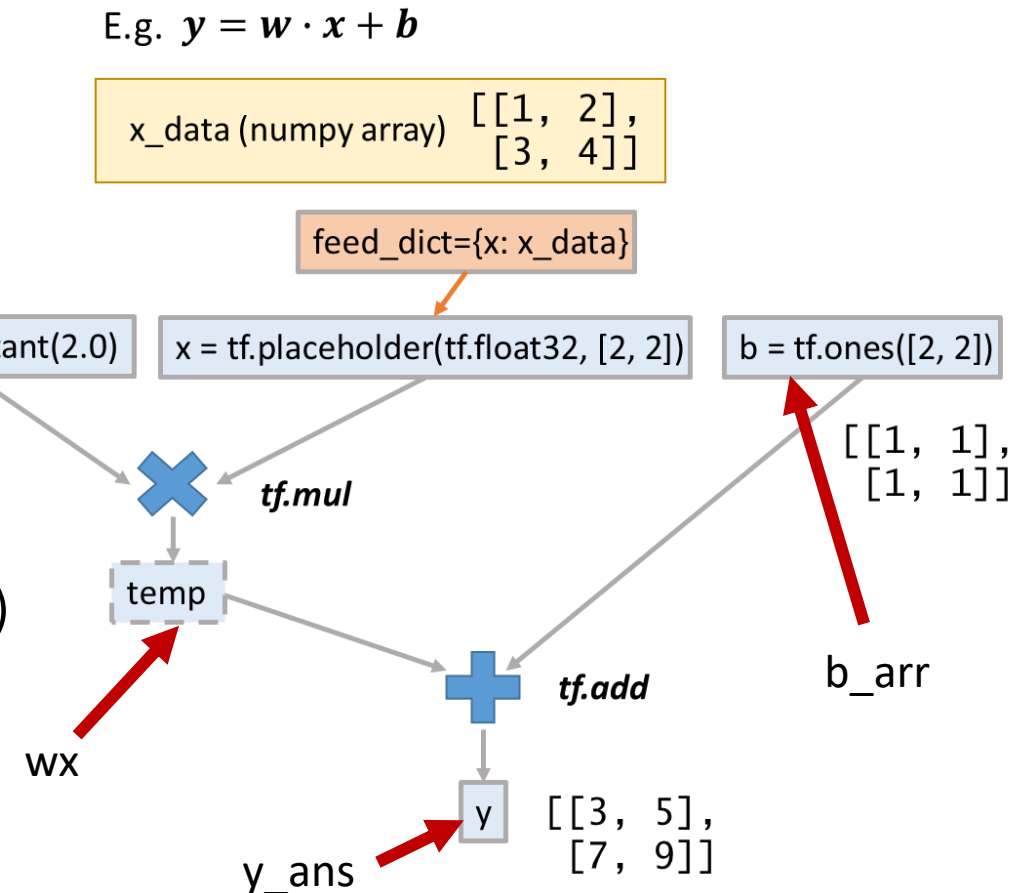
**How to get the result?**

# Run the graph

Usage:

- **`sess.run([<nodes>], <feed_dict>)`**
  - E.g. `b_arr = sess.run(b)`
  - E.g. `wx = sess.run(tf.mul(w, x), feed_dict={x: x_data})`
  - E.g. `y_ans = sess.run(y, feed_dict={x: x_data})`
- **`<tensor>.eval(session=sess)`**
  - E.g. `b_arr = b.eval(session=sess)`
  - E.g. `wx = tf.mul(w, x).eval(feed_dict={x: x_data}, session=sess)`
  - E.g. `y_ans = y.eval(feed_dict={x: x_data}, session=sess)`

***Only run the graph before the node you designate!***





# References

- [Stanford CS224d: TensorFlow Tutorial](#)
- [Stanford CS231n: Deep Learning Software](#)
- [TensorFlow docs](#)

# Environment Setup

# Environment Setup

**# install miniconda (python2.7)**

**# if you don't have wget, you can directly go to the website to download the script**

wget [https://repo.continuum.io/miniconda/Miniconda2-latest-MacOSX-x86\\_64.sh](https://repo.continuum.io/miniconda/Miniconda2-latest-MacOSX-x86_64.sh)

bash Miniconda2-latest-MacOSX-x86\_64.sh

**# append the following line to ~/.bashrc if it is not done automatically**

export PATH="path/to/anaconda2/bin":\$PATH

**# check conda installation**

conda list

# Environment Setup

## **# create new environment**

conda create -n tensorflow

## **# activate the environment**

source activate tensorflow

## **# install required package in the environment**

pip install opencv

pip install matplotlib

pip install --ignore-installed --upgrade

<https://storage.googleapis.com/tensorflow/mac/cpu/tensorflow-1.2.0-py2-none-any.whl>

pip install jupyter notebook


# TensorFlow Sample codes

# Linear Regression

E.g.  $y = W \cdot x + b$

[[2, 6],  
[1, 2],  
[4, 5],  
[6, 8]]

[[43],  
[20],  
[44],  
[65]]



feed\_dict={x: data, y: label}

x = tf.placeholder(tf.float32, [batch\_size, data\_dim])

y = tf.placeholder(tf.float32, [batch\_size, data\_dim])

W = tf.Variable(tf.random\_uniform([data\_dim, 1], -1, 1))

b = tf.Variable(tf.random\_uniform([1], -1, 1))

*Variables are optimized  
during training*

When testing, you only  
run to **y\_pred** or **loss**

y\_pred = tf.add(tf.matmul(x, W), b)

loss = tf.reduce\_mean(tf.square(y - y\_pred))

$$J(W, b) = \frac{1}{N} \sum_{i=1}^N (y - y_{pred})^2$$

Testing  
Training

During training, you need  
to run the **optimizer**

opt = tf.train.AdamOptimizer(learning\_rate=1).minimize(loss)

```
In [2]: import tensorflow as tf
import numpy as np
```

## Linear Regression

$$y = W \cdot x + b$$

Given some data points and their labels, we can learn the parameters (W and b) of the model by reducing the loss.

The answer of this model's parameters are:

W\_ans = [[3, 5]]

b\_ans = [7]

```
In [3]: # data & label
data = np.array([[2, 6], [1, 2], [4, 5], [6, 8]])
label = np.array([[43], [20], [44], [65]])
```

## Build the graph (define your model)

```
In [4]: [batch_size, data_dim] = data.shape
# reserve place for x and y by placeholder
x = tf.placeholder(tf.float32, [batch_size, data_dim])
y = tf.placeholder(tf.float32, [batch_size, 1])

# W and b are random initialized
W = tf.Variable(tf.random_uniform([data_dim, 1], -1, 1))
b = tf.Variable(tf.random_uniform([1], -1, 1))

# y = w*x+b
y_pred = tf.add(tf.matmul(x, W), b)


# compute the loss
loss = tf.reduce_mean(tf.square(y-y_pred))

# declare an optimizer
opt = tf.train.AdamOptimizer(learning_rate=1).minimize(loss)
```

## Training and Testing

```
In [37]: # create a session
sess = tf.Session()
# initialize variables
```

# Mnist classification using Neural Network



$$\text{cross-entropy loss} = - \sum_{i=1}^N y_i * \log(\hat{y}_i)$$

y is one hot encoding of the correct class



# Neural Network classification in MNIST digits dataset

```
In [2]: import tensorflow as tf
        from tensorflow.examples.tutorials.mnist import input_data
        import numpy as np
```

```
In [3]: # Define training parameter
        max_iter = 10000
        batch_size = 100

        # get MNIST data
        mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
```

```
Successfully downloaded train-images-idx3-ubyte.gz 9912422 bytes.
Extracting MNIST_data/train-images-idx3-ubyte.gz
Successfully downloaded train-labels-idx1-ubyte.gz 28881 bytes.
Extracting MNIST_data/train-labels-idx1-ubyte.gz
Successfully downloaded t10k-images-idx3-ubyte.gz 1648877 bytes.
Extracting MNIST_data/t10k-images-idx3-ubyte.gz
Successfully downloaded t10k-labels-idx1-ubyte.gz 4542 bytes.
Extracting MNIST_data/t10k-labels-idx1-ubyte.gz
```

## Function for Neural Network

```
In [4]: def add_layer(inputs, in_size, out_size, activation_function=None):
        # add one more layer and return the output of this layer
        Weights = tf.Variable(tf.random_normal([in_size, out_size]))
        biases = tf.Variable(tf.zeros([1, out_size]) + 0.1)
        Wx_plus_b = tf.add(tf.matmul(inputs, Weights), biases)
        # apply dropout and the given activation function
        if activation_function is None:
            outputs = tf.nn.dropout(Wx_plus_b, keep_prob)
        else:
            outputs = tf.nn.dropout(activation_function(Wx_plus_b), keep_prob)
        return outputs
```

```
In [5]: def compute_accuracy(x, y):
        y_pre = sess.run(prediction, feed_dict={xs: x, keep_prob: 1.})
        correct_prediction = tf.equal(tf.argmax(y_pre, 1), tf.argmax(y, 1))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
        result = sess.run(accuracy, feed_dict={xs: x, ys: y, keep_prob: 1.})
        return result
```

CNN sample codes

# Convolutional Neural Networks are just Neural Networks BUT:

## 1. Local connectivity





image: 32x32x3 volume


**before:** full connectivity: 32x32x3 weights

**now:** one neuron will connect to, e.g. 5x5x3 chunk and only have 5x5x3 weights.

note that connectivity is:

- local in space (5x5 inside 32x32)
- but full in depth (all 3 depth channels)


Replicate this column of hidden neurons across space, with some **stride**.



7x7 input

assume 3x3 connectivity, stride 1


Replicate this column of hidden neurons across space, with some **stride**.



7x7 input


assume 3x3 connectivity, stride 1

Replicate this column of hidden neurons across space, with some **stride**.



7x7 input  
assume 3x3 connectivity, stride 1


Replicate this column of hidden neurons across space, with some **stride**.



7x7 input  
assume 3x3 connectivity, stride 1




Replicate this column of hidden neurons across space, with some **stride**.



7x7 input  
assume 3x3 connectivity, stride 1




Replicate this column of hidden neurons across space, with some **stride**.



7x7 input  
assume 3x3 connectivity, stride 1  
=> **5x5 output**

Replicate this column of hidden neurons across space, with some **stride**.




7x7 input

assume 3x3 connectivity, stride 1

=> **5x5 output**

what about stride 2?

Replicate this column of hidden neurons across space, with some **stride**.




7x7 input

assume 3x3 connectivity, stride 1

=> **5x5 output**

what about stride 2?

Replicate this column of hidden neurons across space, with some **stride**.




7x7 input

assume 3x3 connectivity, stride 1

=> **5x5 output**

what about stride 2?

Replicate this column of hidden neurons across space, with some **stride**.




7x7 input

assume 3x3 connectivity, stride 1

=> **5x5 output**

what about stride 2?

=> **3x3 output**



Output size:  
 $(N - F) / \text{stride} + 1$

e.g.  $N = 7, F = 3$ :

stride 1  $\Rightarrow (7 - 3) / 1 + 1 = 5$

stride 2  $\Rightarrow (7 - 3) / 2 + 1 = 3$

stride 3  $\Rightarrow (7 - 3) / 3 + 1 = \dots \backslash$




Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5**, **stride 1**

Number of neurons: **5**



Output volume:  $(32 - 5) / 1 + 1 = 28$ , so: **28x28x5**

How many weights for each of the 28x28x5  
neurons? **5x5x3 = 75**

# In practice: Common to zero pad the border

(in each channel)

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

neuron with receptive field 3x3, stride 1

pad with 1 pixel border => what is the output?

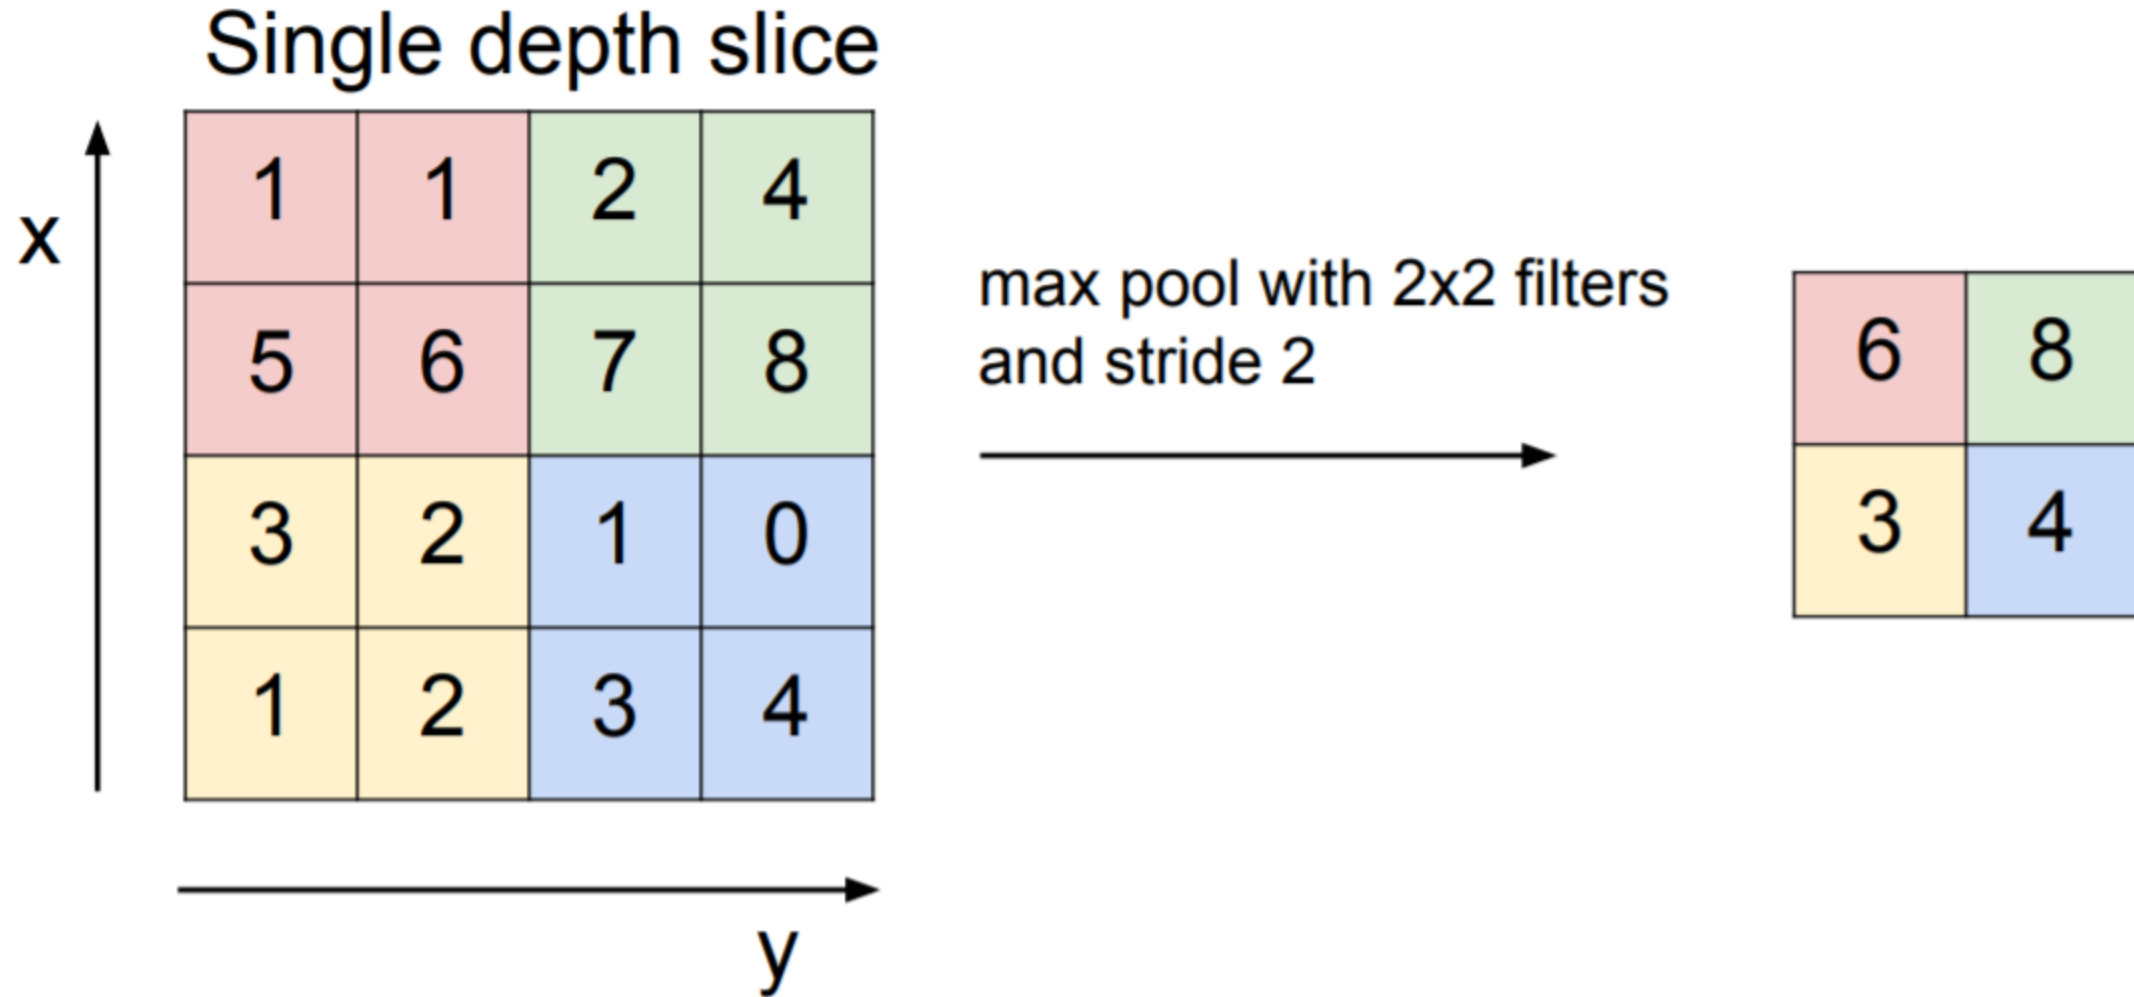
7x7 => preserved size!

in general, common to see stride 1, size  $F$ , and zero-padding with  $(F-1)/2$ .

(Will preserve input size spatially)

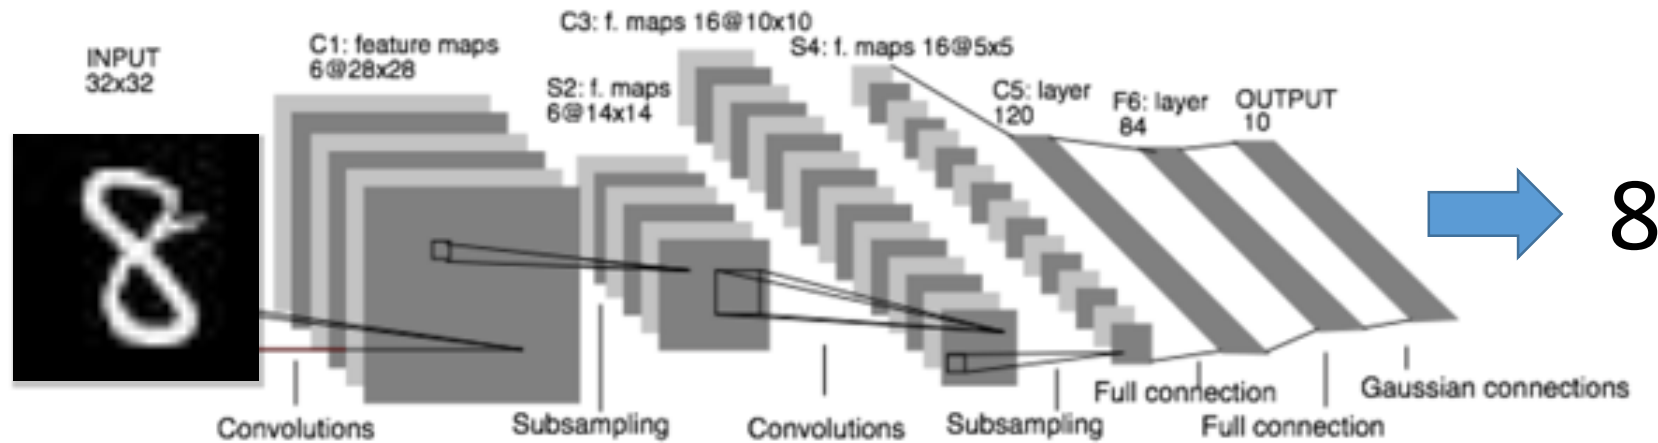


# MAX POOLING



# Mnist classification using LeNet


## Convolutional Networks: 1989



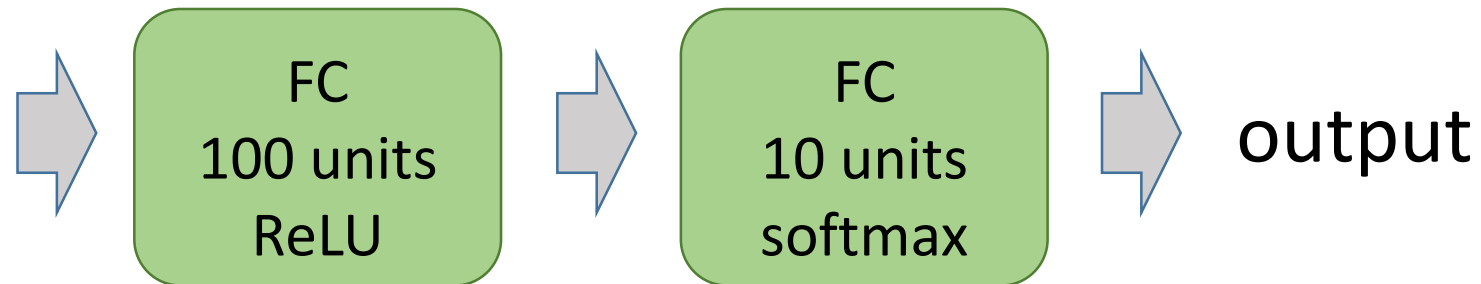
LeNet: a layered model composed of convolution and subsampling operations followed by a holistic representation and ultimately a classifier for handwritten digits. [ LeNet ]

# CNN model for mnist

## PART I




## PART II



RNN sample codes

# Mnist classification using RNN



# References

- [Stanford CS231: Convolutional Neural Networks for Visual Recognition](#)